



# Reinforcement Learning for Goal Recognition

Gabriel Waengertner Henrique

# Introduction



- Reinforcement Learning
  - Capable of learning robust policies for environment interaction
  - Policies can be made to follow any desired behavior
- Goal Recognition is an important ability for interaction
  - Intelligent agents implicitly recognize objectives given a behavior
  - Observations can be noisy

# Introduction




**Can we learn policies using reinforcement learning  
in order to recognize goals?**

# Reinforcement Learning

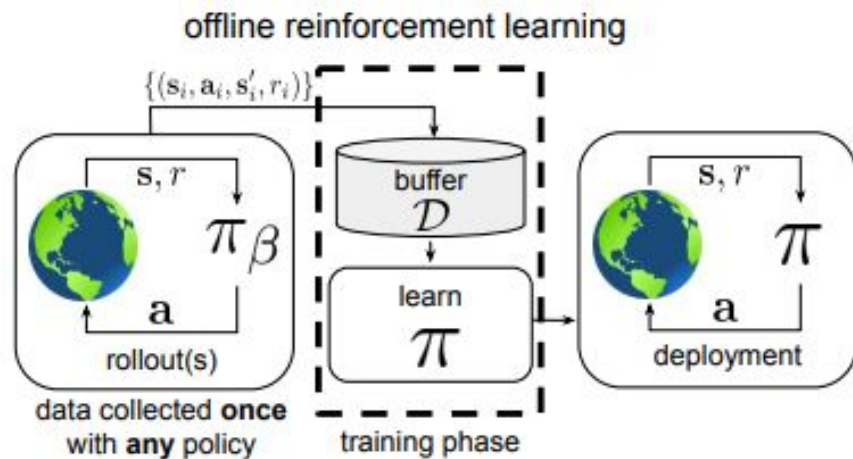
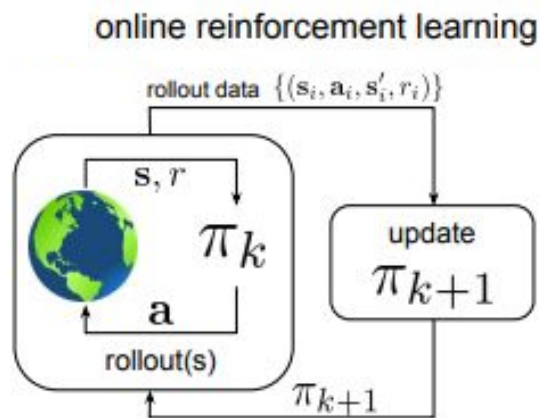


- A Reinforcement Learning (RL) problem consists of an MDP  $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ 
  - $\mathcal{S}$  is a finite set of states
  - $\mathcal{A}$  is a finite set of actions
  - $\mathcal{P}$  is a transition function
  - $\mathcal{R}$  is a reward function
  - $\gamma \in [0, 1]$  is a discounting factor
- Learn a policy  $\pi$  that maximizes the expected discounted reward

# Reinforcement Learning



# Reinforcement Learning



# Goal Recognition



- A Goal Recognition (GR) problem is a tuple  $P_G = \langle \Xi, \mathcal{I}, \mathcal{G}, O \rangle$ , where
  - $\Xi$  is the domain definition
  - $\mathcal{I}$  is an initial state
  - $\mathcal{G}$  is a finite set of candidate goals
  - $O$  is a sequence of observations  $\langle o_1, o_2, \dots, o_n \rangle$
- The solution to a GR problem is a goal  $g \in \mathcal{G}$  that is most consistent with observations  $O$

# Goal Recognition



- Observations
  - Noisy
  - Missing
  - Different representations
- Datasets available<sup>1</sup>

<sup>1</sup> <https://github.com/pucrs-automated-planning/goal-plan-recognition-dataset>



# Proposal



- Use Reinforcement Learning as a solution to a Goal Recognition Problem
- Initial steps on a larger study
  - Learn domain definitions using Reinforcement Learning

# First Steps



- Train a tabular RL method to learn a different policy for each goal
- Find the most likely goal
  - Extract a trajectory using an off-the-shelf planner
  - Measure distance between policies when following a trajectory
  - KL Divergence
- Toy domains
  - Blocksworld

# Schedule



Week 1 - Develop tabular Reinforcement Learning methods to learn policies in small state space domains

Week 2 - Start testing Inverse Reinforcement Learning as a solution to Goal Recognition

Week 3 - Use functions approximators, starting with linear ones

Week 4 - Expand to larger state space domains

Week 5 - Compare with other works and write final paper with results



# Thank you!

Reinforcement Learning for Goal Recognition  
Gabriel Waengertner Henrique